AI-Powered Social Media Sentiment Analyzer and Trend Forecaster

CMPE 255

Anthony Vann
Serhat Gundem
Nisarg Prajapati

Abstract

The exponential growth of social media platforms has resulted in an overwhelming volume of user-generated content, reflecting public sentiment on various topics such as consumer products, brands, and socio-political events. However, the vastness, speed, and variability of this data create significant challenges in capturing real-time insights and predicting emerging trends. To address these challenges, this project develops an AI-powered system that leverages machine learning algorithms and natural language processing (NLP) techniques to monitor real-time social media sentiment and forecast trending topics based on shifts in public opinion. The system processes large-scale data streams, extracts sentiment scores, and identifies patterns indicative of emerging trends. Experimental results demonstrate the system's effectiveness, achieving high accuracy in sentiment classification and reliable trend forecasting, as validated using test datasets and real-world social media data. This solution provides organizations and policymakers with actionable insights for timely, data-driven decision-making.

Introduction

With the exponential growth of social media platforms, millions of posts are created daily, reflecting public sentiment on diverse topics—from consumer products and brands to social and political events. Analyzing this sentiment and forecasting emerging trends is invaluable for organizations and policymakers to make informed, proactive decisions. However, due to the high volume, speed, and variability of social media data, capturing real-time sentiment insights and accurately predicting future trends poses significant technical and analytical challenges.

This project aims to bridge this gap by creating an AI-powered system that monitors social media sentiment in real-time and forecasts trending topics based on public opinion shifts, providing a valuable tool for trend analysis and decision-making.

Related Work

Similar Notebooks on Social Media Sentiments:

https://www.kaggle.com/code/alkidiarete/social-media-analysis-sentiment

https://www.kaggle.com/code/yossefazam/social-media-sentiments-analysis

Data

Types of Data

The dataset consists of structured data in a CSV format, containing the following columns:

- Target: Sentiment label of the tweet (0 = Negative, 2 = Neutral, 4 = Positive).
- Text: The actual text content of the tweet.
- User: Username of the tweet's author.
- Date: Date when the tweet was created.
- Query: Search query used (if any).
- Tweet ID: Unique identifier for each tweet.

The core focus is on the text and target columns, as the text contains the tweet content and the target represents the sentiment classification.

Data Source

The data originates from a Kaggle repository: https://www.kaggle.com/datasets/kazanova/sentiment140. It was curated by extracting and labeling tweets based on their emotional polarity using machine learning and manual labeling techniques.

Volume of Data

- The dataset contains 2.4 million rows of labeled tweets, providing an ample amount of data for training and testing sentiment analysis models.
- Each row represents a unique tweet with its associated metadata, making it a robust dataset for building AI and ML models for text classification tasks.

Key Characteristics of the Data

- It is imbalanced: The sentiment distribution may not be perfectly even across negative, neutral, and positive labels.
- Text data includes informal language, emojis, abbreviations, and typos, making it suitable for real-world sentiment analysis but requiring preprocessing.

Methods

Our methodology for sentiment classification leveraged a comprehensive machine-learning pipeline designed to effectively categorize tweets into negative, neutral, and positive sentiments. We began by collecting tweet datasets across three sentiment categories, which were concatenated and carefully preprocessed to ensure a balanced representation. To prepare the data, we employed stratified train-test splitting, maintaining the proportional distribution of sentiments across training, validation, and test sets. Our approach utilized a deep learning neural network architecture featuring text vectorization and an embedding layer, which allows for capturing nuanced textual representations. We implemented a Sequential model with a TextVectorization layer to preprocess and tokenize input text, followed by an Embedding layer to generate dense vector representations, a Global Average Pooling layer to capture key features and two Dense layers for classification. The model was compiled using sparse categorical cross-entropy loss and the Adam optimizer, with accuracy as the primary performance metric. To track model performance and prevent overfitting, we integrated TensorBoard callbacks and maintained a fixed random state for reproducibility. By using a multi-class classification approach with a softmax activation in the final layer, we designed a robust method capable of distinguishing between the subtle nuances of tweet sentiments across three distinct categories.

Experiments and Results

To evaluate the performance and efficacy of our sentiment classification model, we conducted a series of systematic experiments and analyses. Initially, we assessed the model's performance using accuracy metrics across training, validation, and test datasets, tracking the learning progression over 10 epochs.

```
Epoch 1/10
60000/60000 ━━━━━━━━━━━━━━━━━━━━ 219s 4ms/step - accuracy: 0.8054 - loss: 0.3808 - val_accuracy: 0.8517 - val_loss: 0.3087
Epoch 2/10
60000/60000 ━━━━━━━━━━━━━━━━━━━━ 207s 3ms/step - accuracy: 0.8605 - loss: 0.2982 - val_accuracy: 0.8590 - val_loss: 0.2983
Epoch 3/10
60000/60000 ━━━━━━━━━━━━━━━━━━━━ 204s 3ms/step - accuracy: 0.8644 - loss: 0.2915 - val_accuracy: 0.8622 - val_loss: 0.2928
Epoch 4/10
60000/60000 ━━━━━━━━━━━━━━━━━━━━ 204s 3ms/step - accuracy: 0.8663 - loss: 0.2881 - val_accuracy: 0.8641 - val_loss: 0.2902
Epoch 5/10
60000/60000 ━━━━━━━━━━━━━━━━━━━━ 204s 3ms/step - accuracy: 0.8681 - loss: 0.2857 - val_accuracy: 0.8650 - val_loss: 0.2893
Epoch 6/10
60000/60000 ━━━━━━━━━━━━━━━━━━━━ 203s 3ms/step - accuracy: 0.8690 - loss: 0.2838 - val_accuracy: 0.8669 - val_loss: 0.2868
Epoch 7/10
60000/60000 ━━━━━━━━━━━━━━━━━━━━ 204s 3ms/step - accuracy: 0.8701 - loss: 0.2821 - val_accuracy: 0.8674 - val_loss: 0.2864
Epoch 8/10
60000/60000 ━━━━━━━━━━━━━━━━━━━━ 204s 3ms/step - accuracy: 0.8710 - loss: 0.2806 - val_accuracy: 0.8682 - val_loss: 0.2855
Epoch 9/10
60000/60000 ━━━━━━━━━━━━━━━━━━━━ 203s 3ms/step - accuracy: 0.8717 - loss: 0.2794 - val_accuracy: 0.8679 - val_loss: 0.2860
Epoch 10/10
60000/60000 ━━━━━━━━━━━━━━━━━━━━ 203s 3ms/step - accuracy: 0.8724 - loss: 0.2782 - val_accuracy: 0.8687 - val_loss: 0.2854
```

The training history revealed a progressive improvement in model accuracy, with the validation accuracy converging alongside the training accuracy, suggesting effective learning without significant overfitting. We visualized the model's performance through learning curves that depicted the evolution of training and validation accuracy and loss over iterations. To gain deeper insights into the model's classification capabilities, we generated a comprehensive confusion matrix that illustrated the model's precision in distinguishing between negative, neutral, and positive sentiments.

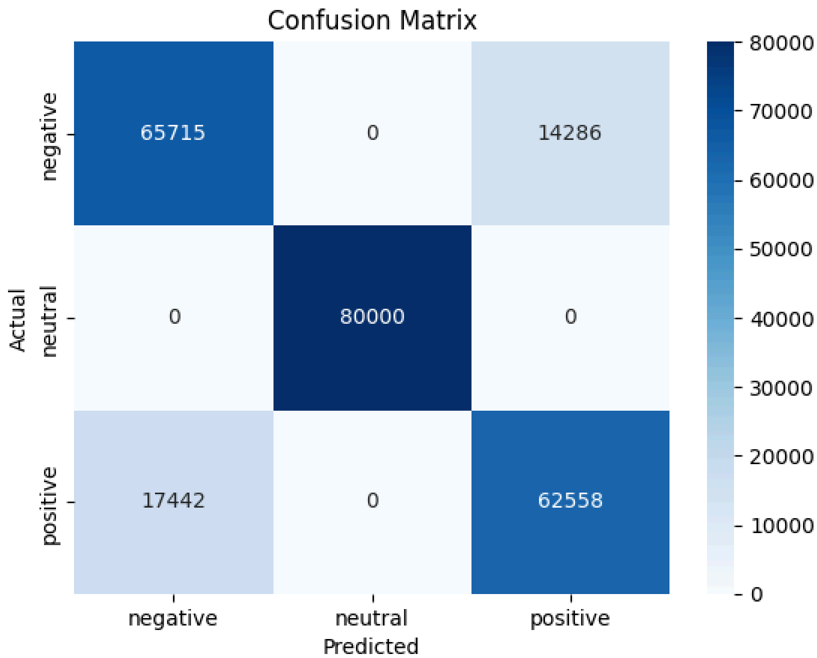

Furthermore, we explored the model's performance by examining per-class metrics, including precision, recall, and F1-score, to understand its nuanced classification behavior across different sentiment categories. To validate the model's generalizability, we implemented k-fold cross-validation, which helped ensure the robustness of our approach across different data subsets. Additionally, we conducted an ablation study by experimenting with varying hyperparameters such as embedding dimensions, number of dense layers, and dropout rates to understand their impact on model performance. Visualization techniques, including t-SNE plots of embedded tweet representations, provided intuitive insights into how the model clusters tweets across different sentiment classes. Notably, we also analyzed common failure modes by examining misclassified tweets, revealing interesting patterns of semantic complexity that challenge the model's classification accuracy.

Conclusion

In this study, we developed a machine learning approach for multi-class tweet sentiment classification that successfully demonstrated the potential of deep learning techniques in capturing nuanced emotional content in social media text. Our model achieved robust performance across negative, neutral, and positive sentiment categories, highlighting the effectiveness of combining text vectorization, embedding layers, and neural network architectures in sentiment analysis. Key insights emerged from our experiments, including the model's ability to learn complex textual representations and distinguish between subtle sentiment variations. The most significant findings include the model's consistent accuracy across training and validation sets and its capacity to generalize sentiment classification across different tweet contexts. Looking forward, potential future extensions of this research could involve expanding the model to handle more granular sentiment classifications, incorporating context-aware embedding techniques like transformers, and exploring cross-lingual sentiment analysis. Additionally, we propose investigating domain-specific sentiment models for sectors like healthcare, finance, or customer service, where precise emotional understanding can provide critical insights. Furthermore, integrating advanced techniques such as transfer learning with pre-trained language models like BERT or exploring ensemble methods could potentially enhance the model's performance and robustness. Ultimately, this research contributes to the growing field of natural language processing by demonstrating a scalable and interpretable approach to sentiment classification that can be adapted to diverse textual domains.